


A defensible machine-learning workflow protects the final test



Do not use test outcomes to choose features, clean data, or tune the model; that leaks answers into evaluation.