


Model selection happens inside training data; the test set is opened once



Every comparison you make on the test set spends part of its independence.

Accuracy alone hides the minority class: a majority-class rule scores 0.76 here and finds no complaint at all.