


Training drives the loss down; the hidden layer buys a curved decision boundary

Binary cross-entropy during 400 full-batch epochs



Two features, two decision boundaries

