


# Training drives the loss down; the hidden layer buys a curved decision boundary

## Binary cross-entropy during 400 full-batch epochs



## Two features, two decision boundaries

